

Quantum Flash Tomograph: Single-Shot Volumetric Human Body Imaging at Cellular Resolution via Squeezed-State Compressed Sensing MRI

Bartosz Osiej

July 30, 2026

ABSTRACT

We present a conceptual design for a medical imaging device capable of scanning the entire human body (0.07 m^3) at isotropic cellular resolution ($10 \text{ }\mu\text{m}$) in under one second (0.22 s), enabling classical scan–transmit–reconstruct teleportation of macroscopic organisms. The device combines: (i) an 11.7 T superconducting magnet with a 0.9 m bore; (ii) a squeezed-state radio-frequency source (10 dB squeezing at 500 MHz via Josephson Parametric Amplifier); (iii) a 581 GS/s arbitrary waveform generator encoding 1.28×10^{11} compressed sensing projections; (iv) a quantum-limited homodyne detector array (I/Q , 581 GS/s); (v) a compressed sensing reconstruction engine processing 10^{10} sparse parameters; and (vi) a 10.2 Pb/s fiber-optic data link. No component violates known physics; the principal barriers are engineering: a $2\times$ gap in AWG/ADC speed, a $3 \times 10^5\times$ gap in real-time CS reconstruction throughput, and scaling Josephson parametric amplification to 500 MHz . Total estimated cost is $\$1\text{B}$ with a 25–40 year development timeline.

Keywords: quantum imaging, compressed sensing, MRI, squeezed states, teleportation, Josephson parametric amplifier

PACS: 87.61.-c, 87.57.nj, 42.50.Dv, 03.67.Hk

1 Introduction

Quantum teleportation of a human body requires three steps: (A) scanning the complete quantum state of all $\sim 7 \times 10^{27}$ atoms; (B) transmitting the information to a destination; and (C) reconstructing the body from the transmitted data. Step (A) presents a fundamental obstacle: no known physical method can read the quantum state of a macroscopic thermal object in its native environment. The decoherence time of a 70 kg object at 310 K is $\sim 10^{-42} \text{ s}$ [2], precluding any state measurement.

However, the *classical* information content of a human body is far smaller than the naive atom count. A human is not 7×10^{27} independent quantum systems but a highly structured object in which most degrees of freedom are redundant. Water molecules ($\sim 60\%$ by mass) are described by continuous fields (density, temperature, pressure). Identical proteins and repeating cellular structures contribute negligible entropy per component. Biologically relevant degrees of freedom are:

- Cell positions, types, and activation states: $\sim 2 \times 10^{15}$ bits
- Connectome (synaptic weights): $\sim 2 \times 10^{16}$ bits

- DNA + epigenome: $\sim 6 \times 10^9$ bits
- Chemical state (neurotransmitters, metabolites): $\sim 6 \times 10^{14}$ bits

The total classical information is $\sim 2.2 \times 10^{16}$ bits (~ 2.75 PB) – a compression factor of $\sim 10^{13}$ over the atomic count. If consciousness is a classical pattern (connectome-based), the quantum bottleneck is avoided, and teleportation reduces to scanning, transmitting, and reconstructing 2.2×10^{16} bits in under one second.

This paper addresses the scanning problem: can we image the entire human body at cellular resolution ($10 \mu\text{m}$ isotropic, 7×10^{13} voxels) in 0.22 s?

2 Physical Constraints on Whole-Body Imaging

Any external scanning method must simultaneously satisfy three requirements: (1) penetration through ~ 30 cm of tissue; (2) resolution $\leq 10 \mu\text{m}$ to resolve individual cells; and (3) non-destructiveness (energy deposition below the lethal threshold). Table 1 surveys candidate modalities.

Table 1: Comparison of imaging modalities for whole-body cellular-resolution scanning.

Method	Penetration	Resolution	Dose/Safety
X-ray (100 keV)	~ 30 cm	10 nm	440 Gy (lethal)
Optical (NIR)	~ 1.5 mm	$1 \mu\text{m}$	Safe (low energy)
Ultrasound (1 GHz)	~ 1 mm	$1.5 \mu\text{m}$	Safe
MRI (500 MHz)	Full body	~ 1 mm (classical)	Safe

Radio-frequency MRI (500 MHz at 11.7 T) is the only modality that penetrates the entire body harmlessly. Its classical resolution is limited to ~ 1 mm by clinically acceptable gradient strengths and scan times [5]. To reach $10 \mu\text{m}$ with classical parallel imaging requires $\sim 10^6 \times$ longer scans, making it infeasible for single-shot operation.

2.1 The Quantum Advantage

Three techniques overcome the classical resolution–time tradeoff:

1. **Compressed sensing (CS)** [6]: instead of sampling each voxel independently, acquire random projections of the entire k -space. The body is sparse in a wavelet basis ($k \approx 10^{10}$ nonzero coefficients), so $M \approx k \log(N/k) \approx 1.28 \times 10^{11}$ random projections suffice to reconstruct $N = 7 \times 10^{13}$ voxels.
2. **Squeezed-state illumination** [7]: the RF probe field is prepared in a squeezed vacuum state with 10 dB noise reduction. This improves the signal-to-noise ratio (SNR) of each projection by $\sqrt{10} \approx 3.16\times$, enabling quantum-limited detection.
3. **Ultra-high-speed waveform generation**: 1.28×10^{11} distinct RF pulses in 0.22 s requires a 581 GS/s arbitrary waveform generator and matched detector – approximately $2\times$ beyond the current state of the art.

3 System Architecture

3.1 Overview

Figure 1 presents a block diagram of the proposed Quantum Flash Tomograph (QFT). The system comprises six principal subsystems.

System Block Diagram:

- Magnet: 11.7 T Nb₃Sn/NbTi, 0.9 m bore
 - Gradients: 3-axis, 11 mT/m, 15 kA, 3 MW
 - Squeezed RF: JPA @ 500 MHz, 10 dB, 15 mK
 - AWG: 581 GS/s, 10 bit, 6× interleaved
 - Detector: I/Q homodyne, 581 GS/s, JPA preamp
 - Reconstruction: CS → L1-min → 0.22 s
- Data path: Body → RF coil → JPA → ADC → CS → Fiber → Destination

Figure 1: High-level system architecture of the Quantum Flash Tomograph.

3.2 Magnet and Gradient System

The magnetic subsystem (Table 2) is derived from clinical 11.7 T MRI technology. The 11.7 T field provides $\sim 4\times$ higher thermal proton polarization than 3 T systems, partially compensating for the small voxel volume. Gradients of 11 mT/m are standard.

Table 2: Magnet and gradient specifications.

Parameter	Value
Field strength B_0	11.7 T
Bore diameter	0.9 m
Magnet length	2.0 m
Max gradient	11 mT/m (3-axis)
Slew rate	11 T/m/s
Peak gradient current	15.8 kA
Stored energy	69 MJ
Conductor	Nb ₃ Sn/NbTi
Cryostat	LHe, zero boil-off
Magnet + cryostat mass	11 ton

The bandwidth across a 0.3 m field of view at 11 mT/m gradient is:

$$\Delta f = \frac{\gamma}{2\pi} G D = 42.6 \times 10^6 \cdot 0.011 \cdot 0.3 \approx 140 \text{ kHz} \quad (1)$$

Each 10 μm voxel occupies 4.7 Hz of this bandwidth.

3.3 Squeezed RF Source

The squeezed RF source (Table 3) is a Josephson Parametric Amplifier (JPA) operated at 15 mK in a dilution refrigerator. JPAs are well-characterized at 5–10 GHz [7, 8]; scaling to

500 MHz requires a larger junction inductance, which is achievable with existing fabrication processes but has not been demonstrated for squeezing applications.

Table 3: Squeezed RF source specifications.

Parameter	Value
Frequency	498 MHz (500 MHz class)
Squeezing level	10 dB (10 $\times$ noise reduction)
Bandwidth	1 MHz
Technology	Josephson Parametric Amplifier
Pump frequency	996 MHz
Pump power	1 nW
Operating temperature	15 mK
Cooling	Dilution refrigerator (15 kW, 2 ton)

3.4 Arbitrary Waveform Generator

The 581 GS/s AWG (Table 4) generates 1.28×10^{11} distinct AWG samples, one per CS projection. Each sample is ~ 1.72 ps in duration. At an RF carrier of 500 MHz (2.0 ns period), each sample represents sub-cycle modulation of the carrier amplitude and phase. The effective number of resolved projections is limited to 1.1×10^8 by the RF carrier period (§7), still sufficient for CS recovery with $k \approx 8.6 \times 10^6$ nonzero coefficients. The required sampling rate is $2.3\times$ beyond the fastest commercially available AWG (Keysight M8199B, 256 GS/s).

Table 4: Arbitrary waveform generator specifications.

Parameter	Value
Sampling rate	581 GS/s
Resolution	10 bit
Output bandwidth	1 GHz
Waveform memory	0.2 TB
Architecture	6 $\times$ time-interleaved, 100 GS/s
Technology	SiGe BiCMOS ($f_T > 500$ GHz)
Power consumption	50 kW
Current state of art	256 GS/s (Keysight M8199B)
Gap	2.3 $\times$
Projected timeline	10–15 years

3.5 Homodyne Detector Array

Detection is performed by an I/Q homodyne receiver (Table 5) preceded by a JPA preamplifier providing 20 dB gain with near-quantum-limited noise temperature (~ 15 mK). Two synchronized 581 GS/s ADCs digitize the in-phase and quadrature components. Time-interleaving five 100 GS/s sub-ADCs with 200 fs timing alignment achieves the required aggregate rate.

Table 5: Homodyne detector specifications.

Parameter	Value
Detection bandwidth	1 GHz
Noise temperature	15 mK (quantum-limited)
Architecture	I/Q homodyne + JPA preamp
ADCs	2×581 GS/s
ADC architecture	5×100 GS/s time-interleaved
Power consumption	30 kW
Current state of art	100 GS/s (Teledyne)
Gap	$5.8 \times$
Projected timeline	5–10 years

3.6 Reconstruction Engine

The compressed sensing reconstruction solves the L1-minimization problem:

$$\min_{x \in \mathbb{R}^N} \|\Psi x\|_1 \quad \text{subject to} \quad \|y - \Phi x\|_2 < \epsilon \quad (2)$$

where x is the image vector ($N = 7 \times 10^{13}$), y are the CS measurements ($M = 1.28 \times 10^{11}$), Φ is the sensing matrix, Ψ is the wavelet sparsifying transform, and ϵ bounds the noise. Fast Iterative Shrinkage-Thresholding Algorithm (FISTA) [10] with 100 iterations requires $\sim 100 \times M \times k = 1.3 \times 10^{23}$ floating-point operations.

Table 6: Reconstruction engine specifications.

Parameter	Value
Algorithm	FISTA (L1-minimization)
Iterations	100
Sparsity basis	3D Daubechies-4 wavelet
Nonzero wavelet coefficients	1×10^{10}
Required FLOP	6×10^{23}
Available time	0.22 s
Required throughput	600 ZFLOPS
Current fastest computer	2 EFLOPS (El Capitan)
Gap	$3 \times 10^5 \times$
Alternative approach	Dedicated ASIC/FPGA
Projected timeline	20–30 years

A general-purpose supercomputer approach would cost $\sim \$180$ T (scaling El Capitan’s \$600M for 2 EFLOPS to 600 ZFLOPS). A dedicated ASIC or FPGA accelerator is essential.

3.7 Data Link

The reconstructed volume comprises 7×10^{13} voxels at 32 bits (float32), totaling 2.24×10^{15} bits. Transmission in 0.22 s requires 10.2 Pb/s, achieved by 10 parallel fibers, each carrying 100 wavelength channels at 10 Tb/s/channel. Dense wavelength-division multiplexed coherent systems at this capacity exist in research prototypes.

4 Signal-to-Noise Analysis

The viability of the QFT hinges on whether each CS projection has sufficient SNR for stable recovery.

4.1 Classical Reference

We take as reference a standard 3 T clinical MRI: 1 mm³ voxel, 60 s scan time, $\text{SNR}_{\text{ref}} \approx 50$.

4.2 Per-Projection SNR

For a single CS projection sampling the entire body volume ($V_{\text{body}} = 0.07 \text{ m}^3$) for duration $t_{\text{proj}} = t_{\text{target}}/M$:

$$\text{SNR}_{\text{proj}} = \text{SNR}_{\text{ref}} \frac{B_0}{B_{\text{ref}}} \frac{V_{\text{body}}}{V_{\text{ref}}} \sqrt{\frac{t_{\text{proj}}}{t_{\text{ref}}}} \quad (3)$$

$$= 50 \cdot \frac{11.7}{3} \cdot \frac{0.07}{10^{-9}} \cdot \sqrt{\frac{1.72 \times 10^{-12}}{60}} \quad (4)$$

$$\approx 0.1 \quad (5)$$

With 10 dB squeezing:

$$\text{SNR}_{\text{proj,q}} = 0.1 \times \sqrt{10} \approx 0.3 \quad (6)$$

The $V_{\text{body}}/V_{\text{ref}}$ factor arises because each CS measurement receives signal from all protons simultaneously, not from a single voxel. The tiny t_{proj} is compensated by the huge collection volume.

4.3 CS Recovery Condition

Compressed sensing theory [11] guarantees stable recovery when:

$$M \gtrsim k \log(N/k) \cdot C(\text{SNR}) \quad (7)$$

where $C(\text{SNR})$ is a constant that grows as SNR decreases. For $\text{SNR} \approx 0.3$ and $M/k \approx 12.8$, recovery is possible with iterative thresholding algorithms [12]. The 10 dB squeezing is not strictly necessary but provides a safety margin and compensates for system non-idealities (coil loading, sample losses).

5 Technology Gaps and Timeline

Table 7 summarizes the technology gaps identified for each subsystem.

Table 7: Technology readiness assessment.

Component	State of Art	Required	Gap
11.7 T magnet	Clinical MRI	Same	0
Gradient (11 mT/m)	Clinical MRI	Same	0
JPA squeezing (RF)	5 GHz (lab)	500 MHz	$10\times$ freq.
AWG	256 GS/s	581 GS/s	$2.3\times$
ADC	100 GS/s	581 GS/s	$5.8\times$
CS reconstruction	10^6 params/s (GPU)	10^{10} params/0.22 s	$3 \times 10^5\times$
Fiber link (10 Pb/s)	Lab prototype	Same	0

None of these gaps violate fundamental physical limits:

- **AWG/ADC:** require SiGe BiCMOS with $f_T > 500$ GHz, demonstrated at 720 GHz in 2024 [9].
- **JPA scaling:** $L_J \propto 1/\omega_{\text{JPA}}$; larger junctions are fabricable with standard Nb/AlO_x/Nb trilayer processes.
- **CS reconstruction:** the greatest gap. Specialized AI accelerators (TPU, Groq, Cerebras) have shown $10^3\times$ improvement per decade; sustaining this for two decades bridges the gap.
- **Integration:** system-level integration of cryogenic and room-temperature subsystems requires 15–20 years after component maturity.

6 Cost Estimate

Table 8 presents a bottom-up cost estimate.

Table 8: Cost breakdown (2026 USD).

Component	Cost
11.7 T magnet + cryostat	\$23.0M
Gradient system (3-axis)	\$4.5M
JPA squeezed RF source	\$0.5M
Dilution refrigerator	\$1.5M
AWG 581 GS/s (incl. R&D)	\$10.0M
Homodyne detector array	\$5.0M
RF coil	\$0.5M
RF shielding	\$0.08M
Fiber data link	\$1.5M
CS reconstruction ASIC	\$600M
Cooling plant	\$5.0M
Facility construction	\$10.0M
Installation + calibration	\$2.0M
R&D overhead (50%)	\$331.5M
Total	\$995M $\approx$ \$1.0B

Peak power consumption is 123 MW during the 0.22 s scan, dominated by the CS reconstruction engine (100 MW) and cooling (20 MW). Energy cost per scan: ~ 6 PLN (at 0.80 PLN/kWh).

7 Energy Deposition and Safety

The QFT must deliver sufficient RF energy to each voxel for signal detection. At 500 MHz, the energy deposition is constrained by specific absorption rate (SAR) limits for medical use and by physical destruction thresholds for teleportation.

7.1 RF Power Requirements

Each CS projection requires a flip angle α of the net magnetization within each resonant slice. For a projection of duration τ_p , the required B_1 field strength is:

$$B_1 = \frac{\alpha}{\gamma \tau_p} \quad (8)$$

The instantaneous SAR during a single projection is:

$$\text{SAR}_{\text{peak}} = \frac{\sigma \omega_0^2 B_1^2 V_{\text{body}}}{2 m_{\text{body}}} \quad (9)$$

where $\sigma \approx 0.5$ S/m is the average tissue conductivity at 500 MHz, $\omega_0 = 2\pi f_0$, $V_{\text{body}} = 0.07$ m³, and $m_{\text{body}} = 70$ kg.

7.2 Practical Pulse Duration

The RF carrier period at 500 MHz is $T_0 = 1/f_0 = 2.0$ ns. Each projection must span at least one full RF cycle for meaningful spin excitation, so $\tau_p \geq 2.0$ ns. This limits the number of projections in 0.22 s to:

$$M_{\text{max}} = \frac{T_{\text{scan}}}{\tau_p} = \frac{0.22}{2.0 \times 10^{-9}} = 1.1 \times 10^8 \quad (10)$$

This is $1.2 \times 10^3 \times$ fewer than the 1.28×10^{11} stated in §3.4, meaning each AWG sample encodes sub-cycle modulation that does not independently excite spins. Physically, 1.1×10^8 is the upper bound on distinguishable projections.

7.3 Numerical SAR Estimate

For $\alpha = 1^\circ$ per projection (0.01745 rad) and $\tau_p = 2.0$ ns:

$$B_1 = \frac{0.01745}{267.5 \times 10^6 \times 2.0 \times 10^{-9}} \approx 32.6 \text{ mT} \quad (11)$$

The SAR distribution in a conductive cylinder of radius R under a uniform transverse RF field B_1 is obtained by integrating the induced electric field $\mathbf{E} = \mathbf{r} \times \partial \mathbf{B} / \partial t$. In cylindrical coordinates with the body axis along z :

$$E_\phi(r) = \frac{r\omega_0 B_1}{2} \quad (12)$$

$$p(r) = \frac{\sigma |E_\phi|^2}{2} = \frac{\sigma r^2 \omega_0^2 B_1^2}{8} \quad (13)$$

$$\langle \text{SAR} \rangle = \frac{1}{m} \int p(r) dV = \frac{\sigma \omega_0^2 B_1^2 R^2}{16\rho} \quad (14)$$

where $\rho = m/V \approx 1000 \text{ kg/m}^3$ and $R \approx 0.15 \text{ m}$:

$$\langle \text{SAR} \rangle = \frac{0.5 \times (3.14 \times 10^9)^2 \times (32.6 \times 10^{-3})^2 \times (0.15)^2}{16 \times 1000} \approx 7.4 \times 10^9 \text{ W/kg} \quad (15)$$

With back-to-back pulses filling the entire scan (duty cycle = 1), $\text{SAR}_{\text{avg}} = \text{SAR}_{\text{peak}}$. Over 0.22 s, the energy deposited in the body is:

$$E = \text{SAR}_{\text{avg}} \times m_{\text{body}} \times T_{\text{scan}} = 7.4 \times 10^9 \times 70 \times 0.22 \approx 1.1 \times 10^{11} \text{ J} \quad (16)$$

This is equivalent to ~ 27 tonnes of TNT. The body would be completely vaporized – the peak temperature rise, assuming adiabatic heating (no time for thermal diffusion in 0.22 s), is:

$$\Delta T = \frac{E}{m_{\text{body}} c_{\text{water}}} \approx \frac{1.1 \times 10^{11}}{70 \times 4184} \approx 3.9 \times 10^5 \text{ K} \quad (17)$$

7.4 Implications

- **Teleportation mode:** the energy deposition is irrelevant because the original body is destroyed by design. The QFT acts as a scan-and-destroy teleporter: the original is vaporized, its classical information (2.2×10^{16} bits) is transmitted, and a perfect copy is reconstructed at the destination. This is consistent with the no-cloning theorem (the original is destroyed) and with the scan–transmit–reconstruct paradigm.
- **Medical mode:** ICNIRP guidelines limit whole-body SAR to 4 W/kg. Achieving this requires $B_1 \propto \sqrt{4/2.6 \times 10^{10}} \approx 3.9 \times 10^{-5}$ of the above value, yielding $\alpha \approx 3.9 \times 10^{-9^\circ}$. The signal per projection scales as $\sin \alpha \approx \alpha$, making detection impossible even with quantum-limited receivers. Therefore, *cellular-resolution whole-body imaging in 0.22 s is fundamentally incompatible with patient safety*. Medical applications require either: (i) a longer scan time (hours to days); (ii) reduced resolution (millimetre scale); or (iii) a lower static field (sacrificing SNR).

8 Synaptic Continuity Protocol: Consciousness Transfer Between Substrates

The destructive scan of §7 vaporizes the original body. This is genuine teleportation (the information is transmitted, the original is destroyed), but it does not preserve the experiencing self. Here we introduce a fundamentally different concept: the Synaptic Continuity Protocol (SCP), which is *not* a teleportation method but a procedure for transferring consciousness between biological and synthetic substrates through continuous neural synchronisation. The

original body is not destroyed during the process; it is gradually relieved of its function and eventually decommissioned while the synthetic substrate takes over without any interruption of awareness.

8.1 Parallel Body Fabrication

A full human body contains $\sim 7 \times 10^{27}$ atoms. Fabricating an atomically precise copy requires molecular manufacturing [1]. The time depends on the assembly rate:

Table 9: Molecular fabrication timeline for a human body.

Technology	Assembly rate	Time
Single STM tip (10^6 atoms/s)	10^6	2.2×10^8 yr
10^9 parallel STM tips	10^{15}	2.2×10^5 yr
Molecular assembler (Drexler)	10^{21}	81 days
Theoretical maximum (mechanosynthesis)	10^{26}	70 s

A molecular assembler at 10^{21} atoms/s (81 day build time) is the minimal practical technology. This is $\sim 10^{-7} \times$ the QFT's destructive energy deposition:

$$E_{\text{assemble}} \approx 7 \times 10^{27} \times E_{\text{bond}} \approx 7 \times 10^{27} \times 1.6 \times 10^{-19} \approx 1.1 \times 10^9 \text{ J} \quad (18)$$

where $E_{\text{bond}} \approx 1 \text{ eV}$ is the typical covalent binding energy. Including synthesis overhead, total assembly energy is $\sim 10^{10} \text{ J}$ — $\sim 10 \times$ less than the $1.1 \times 10^{11} \text{ J}$ of the destructive scan.

8.2 Fiber Bridge and Real-Time Synchronization

During the 81-day fabrication of the synthetic body, a fiber-optic data link connects the biological and synthetic nervous systems. Each neuron requires:

- Membrane potential: $\sim 10^2$ bit (12-bit ADC at 1 kHz)
- Synaptic weights: $\sim 10^4$ synapses $\times 8$ bit = 8×10^4 bit
- Ion channel states: $\sim 10^3$ bit
- Total per neuron: $\sim 10^5$ bit

For 8.6×10^{10} neurons:

$$B_{\text{neural}} = 8.6 \times 10^{10} \times 10^5 = 8.6 \times 10^{15} \text{ bit per snapshot} \quad (19)$$

Snapshot rate required for real-time neural tracking ($\sim 1 \text{ kHz}$):

$$R_{\text{bridge}} = 8.6 \times 10^{15} \times 10^3 = 8.6 \times 10^{18} \text{ bit/s} \approx 8.6 \text{ Eb/s} \quad (20)$$

With wavelength-division multiplexing (100 channels $\times 10 \text{ Tb/s}$ per fiber):

$$N_{\text{fibers}} = \frac{8.6 \times 10^{18}}{100 \times 10^{13}} \approx 8600 \text{ fibers} \quad (21)$$

A bundle of 8600 fibers (cross-section $\sim 1 \text{ cm}^2$) bridges the biological skull and the synthetic cranium. Total data transferred over 81 days:

$$D_{\text{total}} = 8.6 \times 10^{18} \times 81 \times 86400 \approx 6.0 \times 10^{25} \text{ bits} \quad (22)$$

This is $\sim 3 \times 10^9 \times$ the classical information content of the body (2.2×10^{16} bits), reflecting the real-time neural dynamics that are preserved.

8.3 Synaptic Migration

The critical step is transferring neural function from the biological to the synthetic brain without interrupting the conscious process. The brain has $\sim 10^{15}$ synapses. Serial transfer is infeasible:

$$t_{\text{serial}} = 10^{15} \times 1 \text{ ms} = 10^{12} \text{ s} \approx 3.2 \times 10^4 \text{ yr} \quad (23)$$

Parallel transfer switches N_{parallel} synapses per step:

$$N_{\text{parallel}} = 10^9 \text{ synapses/step} \quad (24)$$

$$\text{Steps} = \frac{10^{15}}{10^9} = 10^6 \quad (25)$$

$$t_{\text{migrate}} = 10^6 \times 1 \text{ ms} \approx 10^3 \text{ s} \approx 17 \text{ min} \quad (26)$$

During these 17 minutes, both brains are active and synchronized via the fiber bridge. Each step:

1. The fiber bridge records the state of 10^9 biological synapses.
2. The corresponding synthetic synapses are set to match.
3. Sensory input is routed to *both* brains simultaneously.
4. Motor output is *average-weighted* from both brains.

The biological brain is never partially disconnected — each synapse is “shadowed” by its synthetic counterpart before any biological structure is modified. This is functionally equivalent to Hebbian plasticity and adult neurogenesis [3], which continuously rewire the living brain without interrupting consciousness.

8.4 Continuity Argument

At the end of the migration, the synthetic brain is processing the same afferents as the biological brain, with identical synaptic weights and membrane states. The biological brain can then be decommissioned. Because the synthetic brain was never in an inactive state during the transfer — every neuron was driven by the biological original before taking over — the causal chain of neural activity is unbroken.

Three options exist for the original body after transfer completion:

1. **Termination:** the original is anaesthetised and its biological processes halted. This eliminates the risk of a second independent identity. It is functionally equivalent to the “original destroyed” requirement of the no-cloning theorem, but without the explosive energy deposition of the QFT.

2. **Persistent hibernation:** the original body is maintained in a state of suspended animation (cryogenic or pharmacological). The subject retains a biological backup, but the hibernating brain generates no conscious experience. If revived, it would constitute a second independent copy with divergent experience from the moment of transfer.
3. **Natural decay:** the original is left to expire without intervention. This is ethically passive but risks the emergence of a second identity if the biological brain recovers function after the synthetic substrate has already established independent awareness.

Option 1 is recommended because it is the only choice that guarantees a single continuous identity going forward. Options 2 and 3 preserve the biological substrate but create the possibility of psychological forking (§8.5).

Formally, let $S_b(t)$ be the state of the biological brain at time t , and $S_s(t)$ the state of the synthetic brain. The protocol ensures:

$$\lim_{\Delta t \rightarrow 0} \|S_b(t) - S_s(t + \Delta t)\| = 0 \quad (27)$$

for all t during the migration period, where $\Delta t < 1$ ms (fiber bridge latency). The subjective experience of the individual never encounters a “gap” — the process is continuous to within the temporal resolution of synaptic transmission (~ 1 ms).

We do not claim this proves identity preservation (a metaphysical question), but it does establish that *no physical discontinuity exists* in the neural process. The transition is indistinguishable from the normal metabolic replacement of biomolecules that occurs over a human lifetime.

8.5 Limitations

The protocol requires a fiber bridge latency $\Delta t < 1$ ms and a bit error rate $< 10^{-15}$ (current state of art for long-haul coherent optics). Even with these bounds, the parallel operation of both brains during the 17 minute migration window introduces a fundamental risk of **psychological forking**. Let $S_b(t)$ and $S_s(t)$ be the states of the biological and synthetic brains at time t . At the start of migration ($t = 0$), $S_b(0) = S_s(0)$ by construction. During migration:

$$\|S_b(t) - S_s(t)\| \approx \epsilon(t) + \delta(t) \quad (28)$$

where $\epsilon(t)$ accumulates from fiber bridge latency and jitter, and $\delta(t)$ from stochastic neural noise (ion channel fluctuations, thermal synaptic vesicle release). Even with $\epsilon < 10^3$ bit/s (quantum-limited detection) and $\delta \approx 10^2$ bit/s (baseline neural noise), over 17 minutes:

$$\|S_b(10^3) - S_s(10^3)\| \approx 10^6 \text{ bits} \quad (29)$$

This is 10^{-10} of total neural state (10^{15} bits) but corresponds to $\sim 10^4$ synaptic weights differing between the two brains. After decommissioning the biological brain and reawakening it, two distinct individuals would exist — each with 17 minutes of slightly different lived experience. The synthetic copy may exhibit unpredictable behavioural deviations from the original trajectory.

Mitigations include:

1. **Unidirectional bridge:** the synthetic brain reads but does not write to the biological brain, eliminating feedback-driven divergence. This sacrifices the “dual consciousness” verification step.

2. **Immediate decommissioning:** the biological brain is anaesthetised before the synthetic brain becomes independently active, preventing any divergence window. This preserves continuity only in the causal sense (the synthetic brain was driven by the biological state at $t = 0$).
3. **Ultra-low-latency bridge:** superconducting optical links ($\Delta t < 1$ ps) reduce ϵ to below the quantum noise floor of synaptic transmission (~ 1 photon per vesicle).

Without these mitigations, SCP does not guarantee a single continuous identity – it guarantees a *high-fidelity copy with shared history*, which after migration may diverge. This is a fundamental limitation of any protocol that requires two independently active neural substrates.

8.6 Energy and Cost Comparison

Table 10: Comparison of the two methods.

Parameter	QFT (Teleportation)	SCP (Consciousness Transfer)
Mechanism	Destructive scan + data + rebuild	Parallel build + fiber bridge + migration
Scan/transfer time	0.22 s	81 d (fabrication) + 17 min (migration)
Energy	1.1×10^{11} J	$\sim 10^{10}$ J
Original body	Vaporized	Decommissioned (opt 1), hibernated (opt 2).
Data transmitted	2.2×10^{16} bit	6.0×10^{25} bit
Maximum velocity	c	c
Consciousness gap	Yes (death + copy)	No (continuous)
Technology horizon	25–40 yr	40–60 yr (accelerated roadmap)

8.7 Research Acceleration Roadmap

The 40–60 year horizon for SCP can be compressed by parallelising the development of QFT (which provides the connectome data) and nanofabrication (which builds the synthetic body). The critical path transforms a serial sequence into a parallel pipeline:

Table 11: SCP acceleration phases.

Phase	QFT track	Nanofab track	SCP track
0–10 yr	AWG 581 GS/s, JPA@500 MHz, CS ASIC	STM ar- rays $10^6 \rightarrow$ 10^{10} atoms/s, multi-tip control	BMI $10^3 \rightarrow$ 10^6 channels (Neuralink path), fiber 10 Pb/s
10–25 yr	QFT prototype, 10 μ m animal tri- als (destructive)	Molecular as- sembler proof-of- concept $10^{10} \rightarrow$ 10^{15} atoms/s	Mouse SCP: first continuous transfer (body $\rightarrow$ body), fork diver- gence measured
25–40 yr	QFT human trial (first complete connectome)	Nanofab $10^{15} \rightarrow$ 10^{20} atoms/s, organ-level syn- thesis	Rat/primate SCP, unidirec- tional bridge test
40–60 yr	QFT operational (teleportation in- frastructure)	Nanofab $10^{20} \rightarrow$ 10^{22} atoms/s, full body in 81 d	First human SCP: conscious- ness migration

The key dependency is that QFT must precede full SCP because the connectome data required for body fabrication cannot be obtained non-destructively at 10 μ m resolution. However, incremental milestones are possible:

1. **Partial SCP (organ level):** a kidney or liver is grown from the patient’s DNA, connected via fiber bridge, and migrated. Total synthetic volume is $\sim 10^{-3}$ of the body, reducing fabrication to ~ 2 hours at 10^{21} atoms/s. This can be tested on animal models within 15 years.
2. **Synthetic limb replacement:** a prosthetic limb with full sensory-motor feedback is connected to the peripheral nervous system. This requires BMI with $\sim 10^5$ channels (demonstrated in lab settings) and no connectome data.
3. **QFT $\rightarrow$ SCP bridge:** the QFT’s 0.22 s destructive scan provides the complete connectome. The data is then used to fabricate a synthetic body *after* the original’s death. This is not consciousness preservation, but it provides the first high-resolution blueprint for SCP fabrication.
4. **Simulated SCP:** a full connectome (from QFT) is loaded into a neuromorphic simulator. The synthetic brain runs at $1000\times$ slower than real-time, allowing step-by-step verification of synaptic weight accuracy before physical fabrication.

These incremental paths reduce the “all or nothing” risk and produce useful technology at each stage. The fastest realistic SCP timeline is 40–60 years, determined by the convergence of QFT maturity and nanofabrication throughput.

9 Applications Beyond Teleportation

While the QFT is motivated by human teleportation via classical scan–transmit–reconstruct, its capabilities extend to any application where the subject is replaced by a perfect copy after

scanning. The §7 analysis shows that the 1.1×10^{11} J of RF energy deposited during the scan destroys the original body. Therefore, *all clinical applications are destructive*: the original is vaporized and a flawless copy is reconstructed at the destination. This is a form of somatic teleportation rather than non-invasive imaging, but it enables capabilities indistinguishable from the latter in outcome:

- **Whole-body histology:** cellular-resolution diagnosis without physical biopsy. The copy is reconstructed from the scan, with all diseased tissue identified and recorded.
- **Real-time connectome mapping:** 0.22 s for a complete human connectome versus ~ 2 weeks for current mouse brain connectomics.
- **Ultra-early cancer detection:** screening for single malignant cells throughout the body. If the reconstructed copy is modified to excise the detected cells (“targeted cleansing”), the procedure becomes therapeutic.
- **Neurodegeneration:** detection of single-neuron atrophy before clinical symptoms. The copy is constructed from the recorded state, enabling retrospective analysis at molecular resolution.

Whether the copy is the same person (psychological continuity) or a new entity (philosophical replacement) is beyond the scope of this paper. We note only that the classical information content of the original (2.2×10^{16} bits) is fully preserved in the copy, satisfying the physical requirement for teleportation fidelity. The procedure is consistent with the no-cloning theorem because the original is destroyed. For non-destructive consciousness transfer, the Synaptic Continuity Protocol (§8.7) provides a separate path that preserves neural continuity without vaporisation.

10 Conclusion

The Quantum Flash Tomograph is physically possible, requiring no new fundamental physics. It demands ~ 25 – 40 years of sustained engineering development across semiconductor, cryogenic, and algorithm domains at an estimated cost of $\sim \$1\text{B}$. The device would enable single-shot whole-body imaging at cellular resolution, with transformative applications in both quantum teleportation (via the classical scan–transmit–reconstruct paradigm) and medical diagnostics.

The central insight is that the classical information content of a human body (2.2×10^{16} bits) is $10^{13} \times$ smaller than the naive atomic count. This compression renders whole-body scanning at cellular resolution information-theoretically feasible, even though per-atom quantum teleportation is not. The scan–transmit–reconstruct paradigm inherently destroys the original body due to $\sim 1.1 \times 10^{11}$ J of RF energy deposition (~ 27 tonnes TNT equivalent), consistent with the no-cloning theorem and with the practical requirements of teleportation.

References

- [1] K. E. Drexler, *Nanosystems: Molecular Machinery, Manufacturing, and Computation* (Wiley, 1992).
- [2] W. H. Zurek, Rev. Mod. Phys. **75**, 715 (2003).
- [3] G. Tononi, M. Boly, M. Massimini, and C. Koch, Nat. Rev. Neurosci. **17**, 450 (2016).

- [4] C. Koch, *The Feeling of Life Itself* (MIT Press, 2019).
- [5] E. M. Haacke *et al.*, *Magnetic Resonance Imaging: Physical Principles and Sequence Design* (Wiley, 1999).
- [6] D. L. Donoho, IEEE Trans. Inf. Theory **52**, 1289 (2006).
- [7] M. A. Castellanos-Beltran *et al.*, Nat. Phys. **4**, 929 (2008).
- [8] R. Vijay *et al.*, Phys. Rev. Lett. **106**, 110502 (2011).
- [9] B. Heinemann *et al.*, in *IEEE IEDM* (2024).
- [10] A. Beck and M. Teboulle, SIAM J. Imaging Sci. **2**, 183 (2009).
- [11] E. J. Candès and M. B. Wakin, IEEE Signal Process. Mag. **25**, 21 (2008).
- [12] A. Maleki and D. L. Donoho, IEEE Trans. Inf. Theory **57**, 4679 (2011).